



## KUBE Documentation

### Introduction:

KUBE is a benchmarking and testing framework rather than a benchmarking 'suite', as it is not our objective to gather and define a bunch of applications or software tests along with proper datasets. This latter is the user's responsibility. As a benchmarking and testing framework, our goal is to create a means to drive, centralize, control and organize the benchmarking and testing process in a consistent, flexible and easy to use way. In this manner any benchmark may be added and managed using KUBE.

KUBE is driven by a configuration file that contains rules to be applied in the different stages of the benchmarking/testing process. The configuration file may be specified as an argument, otherwise the default "etc/kube.yaml" file will be used. The user may also develop custom configuration files (for example if you have different machines and environments) and choose to use them as needed.

### Directory structure:

As a means to understand KUBE's process and configuration file, we first advise understanding the directory structure of the framework.

The directory structure is as follows:

```
|--- bench/   ### Contains categorized benchmarks
|           |--- apps/
|           |--- filesystems/
|           |--- networks/
|           |--- synthetics/
|--- bin/     ### The main scripts live here
|--- doc/
|--- etc/     ### Contains the configuration file(s) and the templates used by KUBE
|           |--- kube.yaml -> noor.yaml
|           |--- noor.yaml
|           |--- run.lsf.in
|--- lib/
|           |--- kube ### contains the core engine scripts
|           |--- PyYAML-3.10
|           |--- yaml -> PyYAML-3.10/lib/yaml/
|
|--- data/
|           |--- results/   ### Contains analysis for each run with the same structure as the 'runs' dir below
|           |           |--- apps/
|           |           |--- filesystems/
|           |           |--- networks/
|           |           |--- synthetics/
|           |--- runs/     ### Contains the actual execution directories for every run
|           |           |--- apps/
|           |           |--- filesystems/
```

```

|         |--- networks/
|         |--- synthetics/
|--- rrd/  ### Contains the RRD databases for the metrics in every run
|         |--- apps/
|         |--- filesystems/
|         |--- networks/
|         |--- synthetics/

```

Inside *bench/apps* there is a list of application names exactly as they appear in the configuration file. Each application should have at least one dataset, but may contain more than one. Inside an app dir, the system should be able to find the files **<dataset\_name>.tgz** being **<dataset\_name>** the name of the dataset as it appears in the configuration file.

The rest of items: filesystems, networks and synthetic are very similar but different to the apps because while each app expects to find one or several self contained datasets as a .tgz files; the rest of items expect to find one or several directory entries to each "dataset". Since filesystems, networks and synthetic can be understood as "synthesized" benchmark (or "artificial", "ad-hoc", ...) they are not likely to contain a dataset in an application sense, but rather a full dir structure with binaries and dependencies that need to be compiled and called with certain arguments. For these reasons, these items have an additional tag in the configuration file: "dependencies". This can be a comma separated list of things that is needed to be copied to the run directory in order to be able to create a self contained dir to run it. Of course it is allowed to use wildcards here, so you can write something like:

```
dependencies: bin/* , lib/*
```

if you need to copy all the contents of the bin dir along with the libs

There are more details regarding the configuration file below.

## Usage:

The main script is a Python script named 'kube'. It's been written in python2.7 and can be called with different argument each defining a specific action: **view**, **run**, **refine** and **probe**. And each action may be controlled with options. The following is the help message shown by kube:

```

$ kube -h
usage:
The accepted ways to run KUBE are:
-----
kube [ {-h,--help} ]
kube {-v,--version}
kube <COMMAND> <OPTIONS>

Some <OPTIONS> are specific to some <COMMANDS>:
-----
kube clean  [-d {runs,results}] [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--since SINCE] [--to TO]
kube view   [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS]
kube run    [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--log FILE]
kube refine [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--since SINCE] [--to TO] [--log FILE] [--rrd] [--force]
kube probe  -t DIR [-p] [-m METRIC_NAMES] [--since SINCE] [--to TO]
kube probe  -b BASE_DIR [-p] [--target DIR] [-m METRIC_NAMES] { [--since SINCE] [--to TO] , [--at DATE] } }

Commands allowed in KUBE:
-----
{clean,view,run,refine,probe}

optional arguments:
-----
-h, --help            show this help message and exit
-v, --version          show program's version number and exit
--since SINCE          Option to set the initial time when needed. Default is the
                        first registered date
--to TO                Option to set the end time when needed. Default is 'Now'
--at DATE              Option to select a specific date when needed

```

<code>--log FILE</code>	Print out log information to FILE
<code>--rrd</code>	Used along with the command 'refine' to enable RRD integration. Each metric computed is stored in a RRD database
<code>--force</code>	The 'refine' stage is only done if the configuration section for the current dataset has changed. You may want to force the 'refine' stage for some reason, ie: create rrd database for specific dates that are already refined
<code>-d {run,results}</code>	Option to specify which directory is to be cleaned when the 'clean' command is used
<code>-a APPS</code>	Option used to select which apps benchmark from the config file will be processed. This option may accept a comma-separated list if multiple apps need to be selected. Use 'all' to select everything. It is allowed in the following commands: 'run', 'view', 'refine'
<code>-n NETS</code>	Option used to select which network benchmark from the config file will be processed. This option may accept a comma-separated list if multiple apps need to be selected. Use 'all' to select everything. It is allowed in the following commands: 'run', 'view', 'refine'
<code>-f FILESYS</code>	Option used to select which filesystem benchmark from the config file will be processed. This option may accept a comma-separated list if multiple apps need to be selected. Use 'all' to select everything. It is allowed in the following commands: 'run', 'view', 'refine'
<code>-s SYNTHS</code>	Option used to select which synthetic benchmark from the config file will be processed. This option may accept a comma-separated list if multiple apps need to be selected. Use 'all' to select everything. It is allowed in the following commands: 'run', 'view', 'refine'
<code>-t DIR</code>	Option to be used with the 'probe' command to select the time analysis. The value expected is the path to the dir containing the post-processed runs
<code>-m METRIC_NAMES</code>	For the time and metrics analysis one or more metrics can be selected using this option. A comma-separated list is used for multiple metrics selection
<code>-b BASE_DIR</code>	Option to be used with the 'probe' command to select the box plot metrics graph analysis
<code>--target DIR</code>	For the metrics graph analysis you may want to specify a dir path which contains one or many runs information to compare against the base dir provided with -k option. This target dir could be either a dir containing a single run for a specific date or a dir containing several runs for various dates or a dataset dir.
<code>-p</code>	When used with -t or -b options, it means that the output will be printed out to stdout instead of creating a graph

Let's go over every action and its options:

## Clean

The clean action is used to remove some contents in order, for example, to free some space on disk or eliminate unnecessary information.

The syntax is:

```
$ kube clean [-d {runs,results}] [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--since SINCE] [--to TO]
```

Without arguments, this action will remove everything stored in 'runs' and 'results' wherever they are, according to what is given in the configuration file. But you may choose to delete only the runs dir (-d runs) or the results dir (-d results). Or even more you can just provide a comma separated list of APPS, NETS, FILESYS or SYNTHS to remove. Additionally you can select a date range to remove using SINCE and TO options. By default SINCE is always the first record and TO is always 'now'.

## View

This is useful to check the current configuration for the whole benchmark, or a specific item. No more needed to say :)

```
$ kube view [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS]
```

## Run

Calling 'run' without arguments will run all the configured benchmarks, or a specific item or a list of items. Normally the info is printed out to the standard output but if the user selects the --log option, Then It will be stored in the specified file. This could be useful, for example when running kube in a cron and want to log the info.

```
$ kube run [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--log FILE]
```

This stage creates a run dir for each item with a timestamp. If for some reason the benchmark can't be submitted or run, the directory is marked as 'INVALID'.

## Refine

'Refine' is a core step since, according to kube's philosophy, runs and results are stored separately. It means that the refine stage is necessary to post-process the runs and keep only the necessary data along with the metrics defined in the configuration file. Runs dir can then be safely removed when not needed or kept for specific users needs. The **lifespan** entry in the configuration file defines when the runs dir will be removed and it is checked every time the refine stage is run. Every run whose metrics cannot be computed is considered 'INVALID'.

```
$ kube refine [-a APPS] [-n NETS] [-f FILESYS] [-s SYNTHS] [--since SINCE] [--to TO] [--log FILE] [--rrd] [--force]
```

Again, as in the 'run' stage, you can refine subsets of the benchmark using the options -a, -n, -f and -s and use time filtering with SINCE and TO and log to file with --log option. But here we have two more additional options: --rrd and --force

The option --rrd enables the RRDTTool integration and the computed metrics are stored in a database. For those who are familiar with RRDTTool, this is the configuration used to store all the metrics:

### RRD step: 6h

**DS name:** name of the metric limited to 19 characters

**Heartbeat:** 48 hours

**RRA 1:** keep one sample every 6 hour for 1 month

**RRA 2:** Weekly average for one year

**RRA 3:** Monthly average for 2 years

From the above you can see that we chose 6h as step. The reason for that is that we don't believe you will be running the benchmark every 5 minutes and may be 6h gives a good resolution. Besides, the heartbeat is set to 48 hours which means that you may not run the benchmark for 2 days and still get valid data.

The --force option can be used every time the user wants to force the refine stage to be done. 'Refine' stage is only done in case the system finds a new run which has not been refined yet or if the section in the configuration file that affects the current datasets changes. But for some reason the user may want to force to do the refine stage.

## Probe

This action is used to analyze and interpret the data, normally through a visual representation. The two types of representations allowed are: metrics' time evolution and metrics comparison using box plots.

Let's start with the time analysis. With the 'probe' command, the option to activate the time analysis is '-t'

```
$ kube probe -t DIR [-p] [-m METRIC_NAMES] [--since SINCE] [--to TO]
```

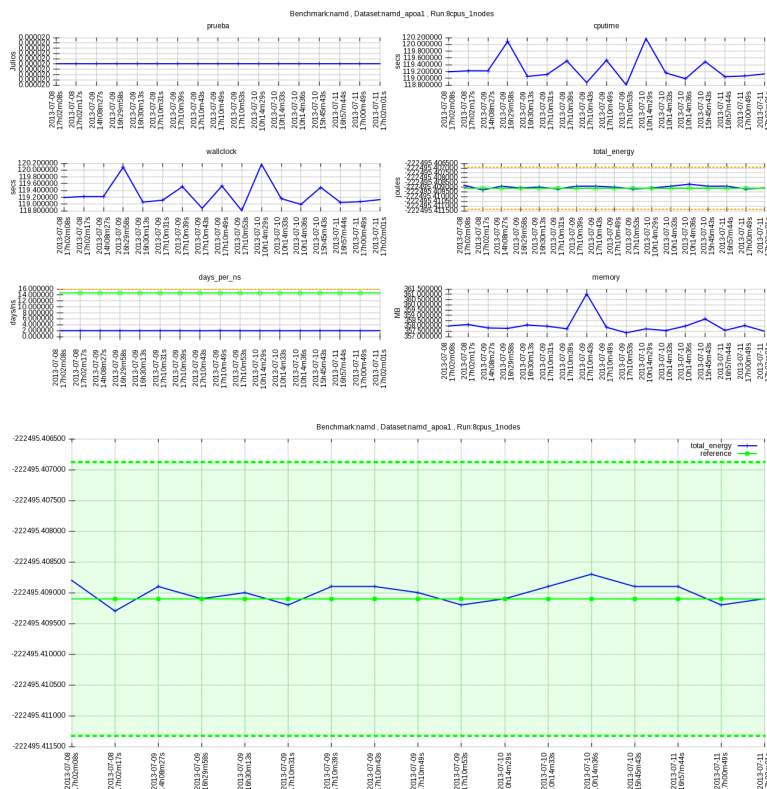
DIR must contain a list of directories named with the Kube's timestamp format. For example, inside DIR you can find subdirs named like:

**2013-07-08T17h02m17s**. The system will search this directories and figure out the metric values to display. User may also want to define just a range of dates to display (a sort of Zoom In effect) using SINCE and TO. The format is the same. For example if you set **--since 2013-07-08T17h** the system will display all the values from results generated after July 8th, 2013 at 5PM until 'now'. Additionally with -m, users can select a list of metrics to display separated by commas. In case the user does not want to display the values but just print them in the screen, they can use the -p option. The columns displayed are:

- 1.- Timestamp
- 2.- Value
- 3.- Reference value
- 4.- Upper valid limit
- 5.- Lower valid limit

The last three are only displayed if they exists.

Here are some examples:



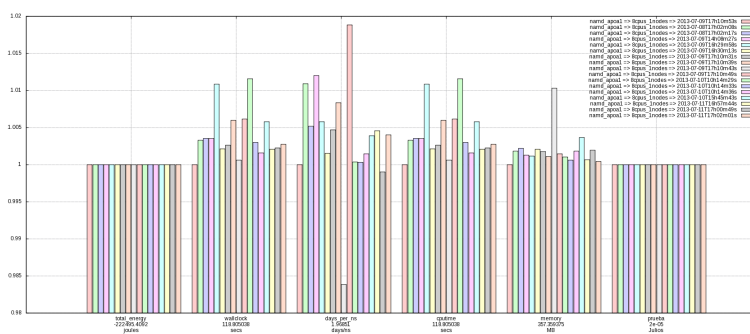
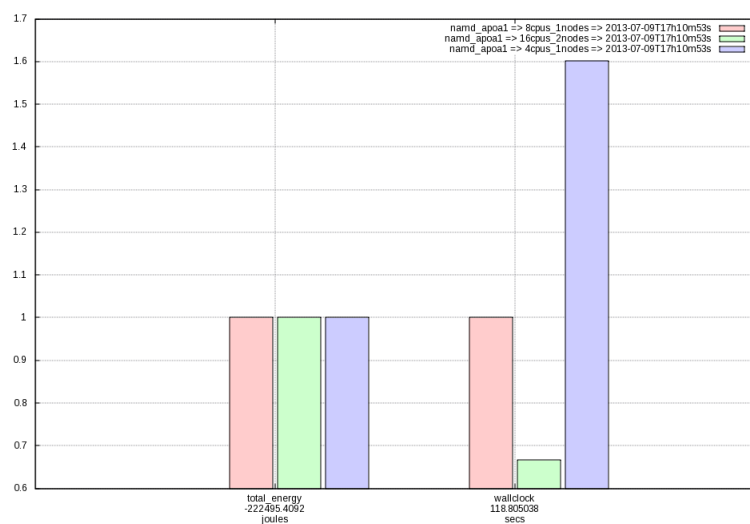
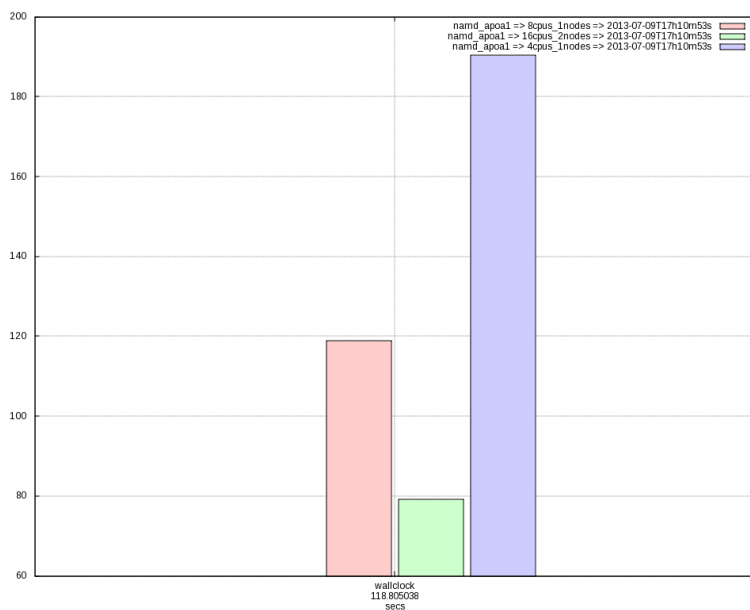
The second way to visualize the metrics is using the box plot diagram.

```
$ kube probe -b BASE_DIR [-p] [--target DIR] [-m METRIC_NAMES] { [--since SINCE] [--to TO] , [--at DATE] }
```

Here BASE\_DIR is the path to the final refined directory, the one with the timestamp name. In this directory Kube will read all the metrics info and will display the information in a box plot or to the screen if the -p option is used. The real added value of the box plot option is to be able to compare the metrics from different runs, usually equivalent runs. For example it makes sense to compare the metrics for the same dataset under different conditions but in the end is up to the user to decide what to compare. To that aim there is the option --target, to select a target DIR to compare with BASE\_DIR. Kube will traverse the target DIR to find the dirs named with the timestamp and read the metrics. The options --since and --to are used to filter the search as you already know. There is an additional option --at to search for a specific date in the target DIR. You could achieve the same with --since and --to but using --at the syntax is more compact and readable.

When displaying a graph, if a single metric is shown, the absolute value is used. But if the user tries to display a combination of metrics into the same plot, Kube will use relative values to avoid scale representations artifacts. So, in this case the values from BASE\_DIR are used as reference and the absolute value is shown in the axis.

Here are some examples:



## Configuration file:

The configuration file is the core component of the framework. The whole process is driven by this file. It is very flexible and there are just a few rules to follow. This flexibility is an advantage and a downside at the same time since it brings you flexibility but the writer is responsible for the mistakes :).

It has been written in YAML because of its simplicity, human readability and easy to parse. So anyone who open the config file could get a quick understanding of the process at a glance.

These are the sections in which the config file is divided:

**Config file fields:****HOME:**

*path:* Full path to the place where KUBE is installed

**RUNS:** An additional field

*path:* Path to the directory where the runs will be stored. If the leading "/" is provided then the absolute path is taken, otherwise the relative path to HOME will be used.

*lifespan:* Defines for how long the runs will be kept. This is an integer value representing days. If 0 keep runs results forever.

**RESULTS:**

*path:* Path to the directory where the results will be stored. If the leading "/" is provided then the absolute path is taken, otherwise the relative path to HOME will be used.

**TOOLS:**

*path:* Path to a special directory where the user can save scripts, binaries or anything that may be use to, for example, process the outputs. If the leading "/" is provided then the absolute path is taken, otherwise the relative path to HOME will be used.

**RDDTOOL:**

*path:* Path to where to store the databases

*root:* This is the path to where the rrdtool is installed. Where to find the binaries

**BATCH:** List of the available batch systems.

*name:* This is a mandatory field and is needed to identify the batch system to use within a benchmark.

*script:* The name of the template used to create the submission script. if the leading "/" is provided then the absolute path is taken, otherwise the relative path to [KUBE][HOME]/etc will be used. Example of the template script used for Noor. Notice that all %<VAR\_NAME>% will be replaced by the specific dataset values at the moment of its creation:

```
#!/bin/bash -l

#BSUB -J %NAME%
#BSUB -o KUBE%NAME%_%NUMPROCS%.out
#BSUB -e KUBE_%NAME%_%NUMPROCS%.err
#BSUB -q %QUEUE%
#BSUB -W %WALLCLOCK%
#BSUB -n %NUMPROCS%
#BSUB %MPIFLAGS%
#BSUB -R span[ptile=%TASKS_PER_NODE%]

module purge
%MODULES%

%LAUNCHER% %LAUNCHER_FLAGS% %EXE% %ARGS%
```

*submit:*

*submittedmsg:* This is the output message (or just part of it) of the submission command containing the JOBID. The output of the submission command will be searched looking for this pattern to get the JOBID and also to determine if the job was submitted correctly to the batch system.

*command:* Submission command

*parameters:* Params to the submission commands if any

If the field 'script' is defined, then the submission command will be:

```
<command> <parameters> <script>
```

A special batch system named **MANUAL** is used in case there isn't any batch system available. This is a mandatory entry that will allow to execute jobs without any batch system. In this case the mandatory tags are '**name**' and '**submit**'. And inside submit, '**command**' and '**parameters**'. For a benchmark using manual launcher ( **MANUAL** batch system ) the submission command to be used would be:

```
<command> <parameters> <exe> <args>
```

Variables %<VAR\_NAME>% specific to the dataset will be replaced at the moment the script is generated. Other %<VAR\_NAME>% variables can be used in the '**submit**' section if they are defined... for example:

```
- name: MANUAL
  hostlistfile: hostnames
  submit:
    command: mpirun
    parameters: -np %NUMPROCS% -machinefile %HOSTSLISTFILE%
```

In this example, '**hostlistfile**' has been added to be replaced in the '**parameters**' section. Notice that %NUMPROCS% will be replaced by the value of the field '**numproc**' of the current run.

The '**batch**' field is mandatory for any dataset, but if the final user doesn't want to use any batch system or manual launcher, for example if you want to run just a sequential test like if you run it from the command line; then he just have to set '**batch**' to '**None**' in the desired dataset:

```
- name: gpfs_iozone3
  active: true
  batch: None
  numprocs: 1
  modules: module load gcc mpi-openmpi
  datasets:
    - name: iozone3_408
      args: -Rab output.wks -g 2G -I
      exe: bin/iozone
      dependencies: bin/*
      active: true
      outputs:
        output: KUBE_%NAME%_%NUMPROCS%.out
        iozone: iozone.tmp
```

**BENCH:** Contains the following fields: *APPS*, *FILESYSTEM*, *SYNTHETIC*, *NETWORKS*. Each of which contains a list of items to be benchmarked.

**APPS:** Contains a list with the applications to be used and its configuration. Mandatory fields for each apps are: '**name**', '**batch**', '**active**', '**dataset**'

*name:* Identifies the app

*batch:* Selects the batch system to be used. Must be a valid name in the BATCH section.

*active:* Indicates whether the app will be run or not

*dataset:* Any active app must have a dataset defined. It is actually a list since one application could run several datasets. Outside the 'dataset' field, user can define whatever tag he wants and it will be used inside each dataset unless the same tag is defined inside any dataset. In this case the tag defined inside the dataset has preference. Note that this tags will be used to match the variables %<VAR\_NAME>% (in capitals) in the submission script (or the submission command) and build up a usable script (or command). So there is no limitation on the tags, except the mandatory ones. The '**numprocs**' tag may contain a comma separated list of values and this tag represents the number of cores to be used

A special mention needs to be done to the **metrics** field since it is a very important component of Kube. This field is a list or array where element defines a metric to be computed for the dataset. Three mandatory fields are necessary and one optional. The mandatory fields are:

*name:* Defines the name of this metric

*units:* Gives more sense to the representation, just defines the units of what is being represented



*command*: Defines the command to be executed in order to compute the metric. It is important to point out here that the expected stream must have the form: <command> | <command> | ....

The optional field is *reference* which defines the range to validate the data computed.

*reference*:

*value*: It is a number or expression. It represents the actual reference value

*tolerance*: May accept the following values: *above* , *below* or *bilateral* and along with the next two fields, define a range of acceptance

*threshold*: Number or "Inf"

*threshold\_type*: May accept *absolute* or *percentage*

Here is an example of an application section:

```
- name: namd
  active: true
  batch: LSF
  # following parameters are dataset dependent...
  # This values will be used unless they are redefined in the "dataset" section
  exe: namd2_mpi
  modules: module load gcc mpi-openmpi libs-extra fftw2/2.1.5-openmpi-1.4.3-sse4.2-sp namd
  metrics:
    - name: wallclock
      units: secs
      command: cat %OUTPUT% | grep WallClock | gawk '{print $2}'
    - name: days/ns
      units: days/ns
      command: cat %OUTPUT% | grep Benchmark | gawk 'END{print $8}'
      reference:
        value: cat %OUTREF% | grep Benchmark | gawk 'END{print $8}'
        tolerance: above
        threshold: inf
        threshold_type: absolute
    - name: cputime
      units: secs
      command: cat %OUTPUT% | grep CPUTime | gawk '{print $4}'
    - name: memory
      units: MB
      command: cat %OUTPUT% | grep Memory | gawk '{print $6}'

  # here goes the specific dataset configuration
  datasets:
    - name: namd_nucleosome146Katoms
      active: false
      args: nucleosome146Katoms.namd > nucleosome.out 2>nucleosome.err
      outputs:
        output: nucleosome.out
        another_output: KUBE_%NAME%_%NUMPROCS%.out

    - name: namd_apoal
      active: true
      numprocs: 2,4
      args: apoal.namd > apoal.out 2>apoal.err
      outputs:
        output: apoal.out
        another_output: KUBE_%NAME%_%NUMPROCS%.out
```

The rest of tags inside '**BENCH**' are :

**FILESYSTEMS:****NETWORKS:****SYNTHETICS:**

They will be explained together since their behavior is similar.

The only differences of these fields compared with the **'APPS'** is that they may contain an additional tag: **'dependencies'** used to copy all the dependencies needed to create a self contained run directory. It allows wildcards and can be a comma separated list of items.

In addition to this, the element contained in **'exe'** will be copied to the execution site if not full path is specified.

For **'APPS'** it is understood that the dataset is self-contained, so no dependencies are allowed.

**Some additional 'rules' to follow in the config file:**

- Only **'outputs'** , **'args'**, **'dependencies'** and **'exe'** fields may contain references to fields as `%<VAR_NAME>%`
- The basic and final element is a *dataset*. Any tag not defined at this level will be searched for in the immediate upper level: **'APPS'**, **'FILESYSTEM'**, ... and if not found there, then the corresponding **'BATCH'** entry is searched.
- if `%NUMPROCS%` is used, one entry per proc will be created.
- The **'metrics'** elements may contain references to the outputs names defined above as well as references to other tags in the config file.
- As a rule of thumb, all paths are relative unless the leading `"/"` is provided; then the absolute path is taken, otherwise the relative path is used.
- Any arbitrary field can be used
- The mandatory fields that must be defined when the dataset level is reached are: **name**, **active**, **args**, **exe**, **output** and **numprocs**